

Relational Algebra and SQL

2.4 and 6.1.

Recall:

Relational Algebra (RA)

• Operations on Relations.

Projection

$\pi_{\langle \text{List Expr} \rangle} R$

list of expressions on the attributes of a relation.

Ex. $R(a,b)$

| a | b |
|---|---|
| 1 | 9 |
| 3 | 3 |

① $\pi_a R$

| a |
|---|
| 1 |
| 3 |

② $\pi_{a+s, -b} R$

| a+s | -b |
|-----|----|
| 6 | -9 |
| 8 | -3 |

name of attributes

③ $\pi_{b,a} R$

| b | a |
|---|---|
| 9 | 1 |
| 3 | 3 |

④ $\pi_{-1,a} R$

| "-1", a |
|---------|
| -1 |
| -1 |

SQL:

select <list expr> from R

① SELECT a FROM R

② SELECT a+5, -b FROM R

③ SELECT b, a FROM R

④ SELECT -1, a FROM R

Name of Relation optional!!

SELECT 3;

| |
|---|
| 3 |
| 3 |

 Creates table of one tuple!!

SELECT 'abc', 5.2
=>

| | |
|-------|-----|
| abc | 5.2 |
| 'abc' | 5.2 |

name of attribute.

Tupler.

The result of SELECT is always a relation

Renaming Relations and their attributes.

Sometimes we need to rename tables or their attributes.

$\rho_{\langle \text{new schema} \rangle} R$

Ex:

$R(a,b)$ $\rho_{S(c,d)} R$
renames $R(a,b)$ to
 $S(c,d)$

ding notation: you can rename during the projection.

If we want to rename the projected expression we can do it:

$\Pi_{a \rightarrow c, b \rightarrow d} R \rightarrow S$
Result schema $S(c,d)$

Ex: ① $\Pi_{a+5 \rightarrow x, -b \rightarrow y} R$

| x | y |
|---|----|
| 6 | -9 |
| 8 | -3 |

SQL.

Given $R(a,b)$ $\rho_{S(c,d)} R$

SELECT a, b FROM R as S(c,d)

or

SELECT a as c, b as d FROM R

①

SELECT a+5 AS x, -b AS y FROM R

In practice it is a good idea to write NATURAL JOINS as JOIN USING.

Eg. $R \bowtie S$

SELECT * FROM

R NATURAL JOIN S

The common attributes are a, b.

It is better to write it as

SELECT * FROM

R JOIN S USING(a,b)

$R \bowtie_{a,b} S$

\nwarrow my notation

For these relations

$R \bowtie S = R \bowtie_{a,b} S$

JOIN USING

SQL provides a special variant of NATURAL JOIN in which we can specify the attribute to join by. Ex.

$R(a,b,c)$ and $S(a,b,d)$

The schema of $R \bowtie S$ is (a,b,c,d)

We can specify a join only on a as follows

SELECT * FROM

R JOIN S USING.

The schema is $(a, R.b, S.b, c, d)$

We write it as.

$R \bowtie_a S$.

SELECTION

$\sigma_p R$

p is a predicate on attributes of R

Expressions:

$<, >, <=, =, >=, <=$

different equal

AND, NOT and many others.

Ex:

| | a | b |
|----------|---|-------------|
| $R(a,b)$ | 3 | 2 |
| | 1 | \emptyset |

p evaluated at each tuple.

① $\sigma_{a>1 \text{ OR } b>1} R$

| a | b |
|---|---|
| 3 | 2 |

SQL.

SELECT * FROM R WHERE p

original attributes of R

Ex:

① SELECT * FROM R
WHERE a > 1 OR b > 1

We can combine Π and σ :

Ex: $\Pi_a \sigma_{a > 1 \text{ OR } b > 1} R$

SELECT a FROM R
WHERE a > 1 OR b > 1

NOT equivalent to.

$\sigma_{a > 1 \text{ OR } b > 1} \Pi_a R$

b is not part of $\Pi_a R$.

Questions

What does this return?

1) $\sigma_{\text{FALSE}} R$

2) $\sigma_{\text{TRUE}} R$

$R \overset{O}{\bowtie} S$
R

| a | b | c |
|---|---|-----|
| 1 | y | 2.5 |
| 2 | ⊥ | 3.1 |

SELECT * FROM
R NATURAL RIGHT JOIN S;

$R \overset{O}{\underset{L}{\bowtie}}_{R.a < S.a} S$

| R.a | R.b | S.a | S.b |
|-----|-----|-----|-----|
| 1 | y | 2 | 3.1 |
| 3 | x | ⊥ | ⊥ |

SELECT * FROM
R LEFT JOIN S ON
(R.a < S.a);

$\left\{ \begin{array}{l} \text{LEFT} \\ \text{RIGHT} \end{array} \right\} \text{ JOIN}$

$R \overset{\circ}{\underset{L}{\bowtie}} S$
 $R \overset{\circ}{\underset{R}{\bowtie}} S$

- Compute non-full join
- Add tuples in LEFT or RIGHT relation, padding other attributes with NULL

Example:

$\therefore R(a,b)$

| a | b |
|---|---|
| 3 | x |
| 1 | y |

$S(a,c)$

| a | c |
|---|-----|
| 2 | 3.1 |
| 1 | 2.5 |

$R \overset{\circ}{\underset{L}{\bowtie}} S$

| a | b | c |
|---|---|-----|
| 1 | y | 2.5 |
| 3 | x | 1 |

SELECT * FROM
 R NATURAL LEFT JOIN S;

Other expressions in predicates.

IN

att IN (List)

Ex:

$a \text{ IN } (3, 2, 5)$

\Rightarrow equivalent to $(a = 3 \text{ or } a = 2 \text{ or } a = 5)$

But we can also use a query:

$a \text{ in } (\pi_c S)$

SQL:

$a \text{ IN } (\text{SELECT } c \text{ FROM } S)$

EXISTS

EXISTS (R) true if R not empty

Ex:

$\text{EXISTS } (\sigma_{ass} R)$

Operations on 2 Relations.

Union \cup

Intersection \cap

Difference (except) $-$

Union Compatible

R and S are "union compatible" iff
 $|attrs(R)| = |attrs(S)|$

and the type of the i -th attribute of S is **type compatible** with the type of the i -th attribute of R.

One type t_1 is type compatible with type t_2 if t_1 can be converted to type t_2 .

$A \cup B$
 $A \cap B$
 $A - B$

} Defined only iff
A & B are
union compatible.

FULL $\left\{ \begin{array}{l} \text{NATURAL JOIN} \quad R \bowtie S \\ \text{THETA JOIN} \quad R \bowtie_P S \end{array} \right.$

- Compute non-full join
- Add tuples in R not in join padded with NULL
- Add tuples in S not in join padded with NULL

$\therefore R(a,b) \quad \begin{array}{c|c} a & b \\ \hline 3 & x \\ 1 & y \end{array} \quad S(a,c) \quad \begin{array}{c|c} a & c \\ \hline 2 & 3.1 \\ 1 & 2.5 \end{array}$

$R \bowtie S$

| a | b | c |
|---|---------|---------|
| 1 | y | 2.5 |
| 3 | x | \perp |
| 2 | \perp | 3.1 |

← Represents NULL in RA

SELECT * FROM R NATURAL FULL JOIN S

$R \bowtie S$
 $R.a < S.a$

| R.a | b | S.a | c |
|---------|---------|---------|---------|
| 1 | y | 2 | 3.1 |
| 3 | x | \perp | \perp |
| \perp | \perp | 1 | 2.5 |

SELECT * FROM R FULL JOIN S
ON (R.a > S.a)

Theta Join

$$R \bowtie_p S = \sigma_p (R \times S)$$

SQL:
`SELECT * FROM
 R JOIN S ON (p);`

Cross Product X

$R \times S$

SQL

`SELECT * FROM R, S;`

NATURAL JOIN

$R \bowtie S$

SQL:

`SELECT * FROM R NATURAL JOIN S`

Theta Join

$$R \bowtie_p S = \sigma_p (R \times S)$$

SQL:

`SELECT * FROM
 R JOIN S ON (p);`

UNION

$$t \in R \cup S \Leftrightarrow t \in R \text{ or } t \in S$$

$$t \in R \cap S \Leftrightarrow t \in R \text{ and } t \in S$$

$$t \in R - S \Leftrightarrow t \in R \text{ and } t \notin S$$

Schema of result is schema of first relation.

Ex:

| R (a, b) | | S (c, d) |
|----------|--|----------|
| a b | | c d |
| 1 a | | 1 e |
| 3 x | | 3 x |
| | | 4 f |

| R ∪ S | R ∩ S | R - S |
|-------|-------|-------|
| a b | a b | a b |
| 1 a | | 1 a |
| 3 x | 3 x | |
| 1 e | | |
| 4 f | | |

| S - R |
|-------|
| c d |
| 1 e |
| 4 f |

SQL

TABLE R { UNION
INTERSECT
EXCEPT } TABLE S
or

<QUERY> { UNION
INTERSECT
EXCEPT } <QUERY>

SELECT a, b FROM R

UNION

SELECT c, d FROM S;

Ex:

R(a, b)

| a | b |
|---|---|
| 1 | x |
| 2 | y |

S(a, c)

| a | c |
|---|----|
| 5 | 8 |
| 2 | 12 |

Common attributes = {a}

$$T = R \bowtie S = \Pi_{a,b,c} \sigma_{R.a=S.a}(R \times S)$$

R x S

| R. a | R. b | S. a | S. c |
|------|------|------|------|
| 1 | x | 5 | 8 |
| 1 | x | 2 | 12 |
| 2 | y | 5 | 8 |
| 2 | y | 2 | 12 |

} R.a = S.a

R ⋈ S

| a | b | c |
|---|---|----|
| 2 | y | 12 |

Natural Join \bowtie

Given relations R and S

C is set of attributes of both S and R
with the same name

• if C is empty.

$$R \bowtie S = R \times S$$

• otherwise

$$\pi_{\text{Attr}(R), \text{Attr}(S) - C} \left(\sigma_{\bigwedge_{a_i \in C} R_{a_i} = S_{a_i}} (R \times S) \right)$$

Do not project
both common
attributes (only
the first).

match tuples
with same value in
common attributes.
conjunction over
all common attributes

NULLS (6.1)

SQL has a special value: NULL
 \Rightarrow unknown.

Example:

- Next year champion of the Stanley Cup.
- Grades of students currently enrolled in this course.

• SQL has special considerations for expressions involving NULL

• SQL Logic 3 valued:

- True
- False
- Unknown

• Any expression involving NULL results into UNKNOWN

IMPORTANT

$$\left. \begin{array}{l} X = \text{NULL} \\ X > \text{NULL} \end{array} \right\} \Rightarrow \text{UNKNOWN.}$$

To test if attr is NULL use
 $X \text{ IS NULL}$

Ex:

$\text{NULL} > 5 \Rightarrow \text{UNKNOWN}$

$X \text{ IS NULL} \Rightarrow \text{True if } X \text{ contains NULL}$

UNKNOWN is NOT true

Ex:

$\text{UNKNOWN OR TRUE} \Rightarrow \text{TRUE}$

$\text{UNKNOWN AND TRUE} \Rightarrow \text{FALSE}$
See exercise !!

Text Matching.

Regular expressions. (Postgres)

$\text{expr} \sim \text{RegExp}$

Ex

$a \sim '^ab'$

attribute a starts with string ab

$a \sim '\backslash.txt\$'$

attribute a end with string .txt

Cross product: X

Given relations R and S.

$(r, t) \in R \times S$ iff $r \in R$ and $s \in S$

$R(a, b)$

| a | b |
|---|---|
| 1 | x |
| 2 | y |

$S(c, d)$

| c | d |
|---|----|
| 5 | 8 |
| 2 | 12 |

$T = R \times S$

| a | b | c | d |
|---|---|---|----|
| 1 | x | 5 | 8 |
| 1 | x | 2 | 12 |
| 2 | y | 5 | 8 |
| 2 | y | 2 | 12 |

What is schema of T?